

# Self-Driving Cars

## Lecture 4 – Reinforcement Learning

**Kumar Bipin**

Robotics, Computer Vision, System Software

BE, MS, PhD (MMMTU, IISc, IIIT-Hyderabad)



# Agenda

**4.1** Markov Decision Processes

**4.2** Bellman Optimality and Q-Learning

**4.3** Deep Q-Learning

# 4.1

## Markov Decision Processes

# Reinforcement Learning

## So far:

- ▶ Supervised learning, lots of expert demonstrations required
- ▶ Use of auxiliary, short-term loss functions
  - ▶ Imitation learning: per-frame loss on action
  - ▶ Direct perception: per-frame loss on affordance indicators

## Now:

- ▶ Learning of models based on the loss that we actually care about, e.g.:
  - ▶ Minimize time to target location
  - ▶ Minimize number of collisions
  - ▶ Minimize risk
  - ▶ Maximize comfort
  - ▶ etc.

# Types of Learning

## **Supervised Learning:**

- ▶ Dataset:  $\{(x_i, y_i)\}$  ( $x_i$  = data,  $y_i$  = label)    Goal: Learn mapping  $x \mapsto y$
- ▶ Examples: Classification, regression, imitation learning, affordance learning, etc.

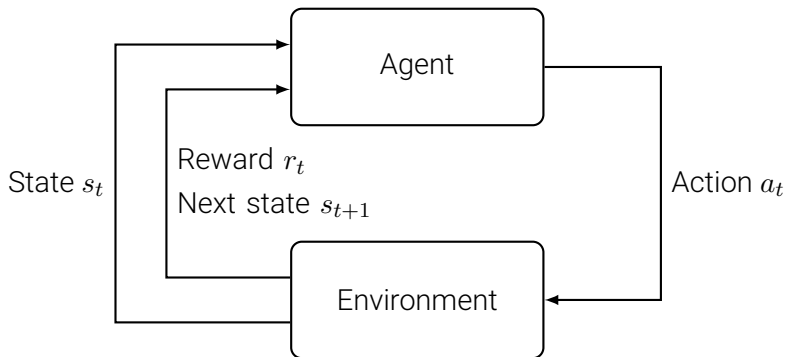
## **Unsupervised Learning:**

- ▶ Dataset:  $\{(x_i)\}$  ( $x_i$  = data)    Goal: Discover structure underlying data
- ▶ Examples: Clustering, dimensionality reduction, feature learning, etc.

## **Reinforcement Learning:**

- ▶ Agent interacting with environment which provides numeric reward signals
- ▶ Goal: Learn how to take actions in order to maximize reward
- ▶ Examples: Learning of manipulation or control tasks (everything that interacts)

# Introduction to Reinforcement Learning

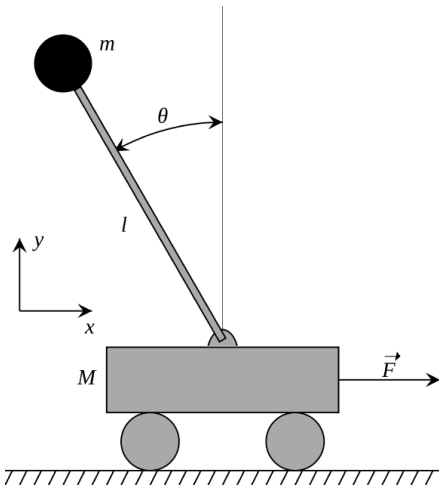


- ▶ Agent observes environment state  $s_t$  at time  $t$
- ▶ Agent sends action  $a_t$  at time  $t$  to the environment
- ▶ Environment returns the reward  $r_t$  and its new state  $s_{t+1}$  to the agent

# Introduction to Reinforcement Learning

- ▶ Goal: Select actions to maximize total future reward
- ▶ Actions may have long term consequences
- ▶ Reward may be delayed, not instantaneous
- ▶ It may be better to sacrifice immediate reward to gain more long-term reward
- ▶ Examples:
  - ▶ Financial investment (may take months to mature)
  - ▶ Refuelling a helicopter (might prevent crash in several hours)
  - ▶ Sacrificing a chess piece (might help winning chances in the future)


## Example: Cart Pole Balancing



- **Objective:** Balance pole on moving cart
- **State:** Angle, angular vel., position, vel.
- **Action:** Horizontal force applied to cart
- **Reward:** 1 if pole is upright at time  $t$



# Example: Robot Locomotion



<http://blog.openai.com/roboschool/>

- ▶ **Objective:** Make robot move forward
- ▶ **State:** Position and angle of joints
- ▶ **Action:** Torques applied on joints
- ▶ **Reward:** 1 if upright & forward moving

# Example: Atari Games



- **Objective:** Maximize game score
- **State:** Raw pixels of screen (210x160)
- **Action:** Left, right, up, down
- **Reward:** Score increase/decrease at  $t$

<http://blog.openai.com/gym-retro/>


## Example: Go



- ▶ **Objective:** Winning the game
- ▶ **State:** Position of all pieces
- ▶ **Action:** Location of next piece
- ▶ **Reward:** 1 if game won, 0 otherwise

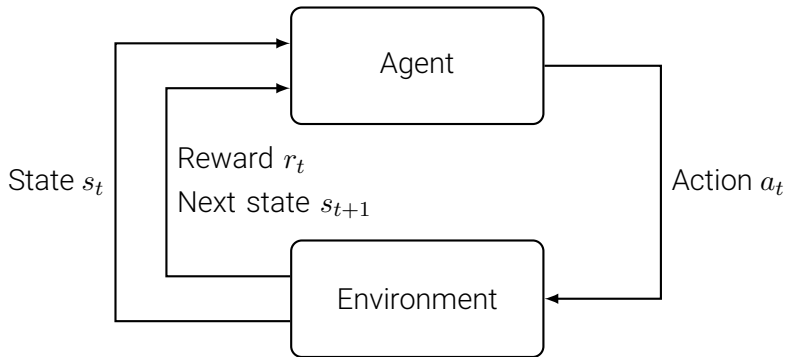
[www.deepmind.com/research/alphago/](http://www.deepmind.com/research/alphago/)

## Example: Self-Driving



- **Objective:** Lane Following
- **State:** Image (96x96)
- **Action:** Acceleration, Steering
- **Reward:** - per frame, + per tile

# Reinforcement Learning: Overview



- How can we mathematically formalize the RL problem?

# Markov Decision Process

**Markov Decision Process (MDP)** models the environment and is defined by the tuple

$$(\mathcal{S}, \mathcal{A}, \mathcal{R}, P, \gamma)$$

with

- ▶  $\mathcal{S}$  : set of possible states
- ▶  $\mathcal{A}$ : set of possible actions
- ▶  $\mathcal{R}(r_t|s_t, a_t)$ : distribution of current reward given (state,action) pair
- ▶  $P(s_{t+1}|s_t, a_t)$ : distribution over next state given (state,action) pair
- ▶  $\gamma$ : discount factor (determines value of future rewards)

Almost all reinforcement learning problems can be formalized as MDPs

# Markov Decision Process

**Markov property:** Current state completely characterizes state of the world

- ▶ A state  $s_t$  is *Markov* if and only if


$$P(s_{t+1}|s_t) = P(s_{t+1}|s_1, \dots, s_t)$$

- ▶ "The future is independent of the past given the present"
- ▶ The state captures all relevant information from the history
- ▶ Once the state is known, the history may be thrown away
- ▶ The state is a sufficient statistics of the future

# Markov Decision Process

## Reinforcement learning loop:

- ▶ At time  $t = 0$ :
  - ▶ Environment samples initial state  $s_0 \sim P(s_0)$
- ▶ Then, for  $t = 0$  until done:
  - ▶ Agent selects action  $a_t$
  - ▶ Environment samples reward  $r_t \sim \mathcal{R}(r_t|s_t, a_t)$
  - ▶ Environment samples next state  $s_{t+1} \sim P(s_{t+1}|s_t, a_t)$
  - ▶ Agent receives reward  $r_t$  and next state  $s_{t+1}$



How do we select an action?



# Policy

A **policy**  $\pi$  is a function from  $\mathcal{S}$  to  $\mathcal{A}$  that specifies what action to take in each state:

- ▶ A policy fully defines the behavior of an agent
- ▶ Deterministic policy:  $a = \pi(s)$
- ▶ Stochastic policy:  $\pi(a|s) = P(a_t = a | s_t = s)$

Remark:

- ▶ MDP policies depend only on the **current state** and not the entire history
- ▶ However, the current state may include past observations

# Policy

How do we learn a policy?

**Imitation Learning:** Learn a policy from **expert demonstrations**

- ▶ Expert demonstrations are provided
- ▶ Supervised learning problem

**Reinforcement Learning:** Learn a policy through **trial-and-error**

- ▶ No expert demonstrations given
- ▶ Agent discovers itself which actions **maximize the expected future reward**
  - ▶ The agent interacts with the environment and obtains reward
  - ▶ The agent discovers good actions and improves its policy  $\pi$

# Exploration vs. Exploitation

How do we discover good actions?

**Answer:** We need to explore the state/action space. Thus RL combines two tasks:

- ▶ **Exploration:** Try a novel action  $a$  in state  $s$ , observe reward  $r_t$ 
  - ▶ Discovers more information about the environment, but sacrifices total reward
  - ▶ Game-playing example: Play a novel experimental move
- ▶ **Exploitation:** Use a previously discovered good action  $a$ 
  - ▶ Exploits known information to maximize reward, but sacrifice unexplored areas
  - ▶ Game-playing example: Play the move you believe is best

**Trade-off:** It is important to explore and exploit simultaneously

# Exploration vs. Exploitation

How to balance exploration and exploitation?

$\epsilon$ -**greedy** exploration algorithm:

- ▶ Try all possible actions with non-zero probability
- ▶ With probability  $\epsilon$  choose an action at random (**exploration**)
- ▶ With probability  $1 - \epsilon$  choose the best action (**exploitation**)
- ▶ Greedy action is defined as best action which was discovered so far
- ▶  $\epsilon$  is large initially and gradually annealed (=reduced) over time



# Value Functions

How good is a state?

The **state-value function**  $V^\pi$  at state  $s_t$  is the expected cumulative discounted reward ( $r_t \sim \mathcal{R}(r_t|s_t, a_t)$ ) when following policy  $\pi$  from state  $s_t$ :

$$V^\pi(s_t) = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t, \pi] = \mathbb{E} \left[ \sum_{k \geq 0} \gamma^k r_{t+k} \middle| s_t, \pi \right]$$

- ▶ The discount factor  $\gamma < 1$  is the value of future rewards at current time  $t$ 
  - ▶ Weights immediate reward higher than future reward  
(e.g.,  $\gamma = \frac{1}{2} \Rightarrow \gamma^k = \frac{1}{1}, \frac{1}{2}, \frac{1}{4}, \frac{1}{8}, \frac{1}{16}, \dots$ )
  - ▶ Determines agent's far/short-sightedness
  - ▶ Avoids infinite returns in cyclic Markov processes

# Value Functions

How good is a state-action pair?

The **action-value function**  $Q^\pi$  at state  $s_t$  and action  $a_t$  is the expected cumulative discounted reward when taking action  $a_t$  in state  $s_t$  and then following the policy  $\pi$ :

$$Q^\pi(s_t, a_t) = \mathbb{E} \left[ \sum_{k \geq 0} \gamma^k r_{t+k} \middle| s_t, a_t, \pi \right]$$

- ▶ The discount factor  $\gamma \in [0, 1]$  is the value of future rewards at current time  $t$ 
  - ▶ Weights immediate reward higher than future reward  
(e.g.,  $\gamma = \frac{1}{2} \Rightarrow \gamma^k = \frac{1}{1}, \frac{1}{2}, \frac{1}{4}, \frac{1}{8}, \frac{1}{16}, \dots$ )
  - ▶ Determines agent's far/short-sightedness
  - ▶ Avoids infinite returns in cyclic Markov processes

# Optimal Value Functions

The **optimal state-value function**  $V^*(s_t)$  is the best  $V^\pi(s_t)$  over all policies  $\pi$ :

$$V^*(s_t) = \max_{\pi} V^\pi(s_t) \qquad V^\pi(s_t) = \mathbb{E} \left[ \sum_{k \geq 0} \gamma^k r_{t+k} \middle| s_t, \pi \right]$$

The **optimal action-value function**  $Q^*(s_t, a_t)$  is the best  $Q^\pi(s_t, a_t)$  over all policies  $\pi$ :

$$Q^*(s_t, a_t) = \max_{\pi} Q^\pi(s_t, a_t) \qquad Q^\pi(s_t, a_t) = \mathbb{E} \left[ \sum_{k \geq 0} \gamma^k r_{t+k} \middle| s_t, a_t, \pi \right]$$

- ▶ The optimal value functions specify the best possible performance in the MDP
- ▶ However, searching over all possible policies  $\pi$  is computationally intractable



# Optimal Policy

If  $Q^*(s_t, a_t)$  would be known, what would be the **optimal policy**?

$$\pi^*(s_t) = \operatorname{argmax}_{a' \in \mathcal{A}} Q^*(s_t, a')$$

- ▶ Unfortunately, searching over all possible policies  $\pi$  is intractable in most cases
- ▶ Thus, determining  $Q^*(s_t, a_t)$  is hard in general (for most interesting problems)
- ▶ Let's have a look at a simple example where the optimal policy is easy to compute

# A Simple Grid World Example

**actions** = {

1. right →

2. left ←

3. up ↑

4. down ↓

}

**states**


★			
			★

**reward:**  $r = -1$  for  
each transition


**Objective:** Reach one of terminal states (marked with '★') in least number of actions

- Penalty (negative reward) given for every transition made

# A Simple Grid World Example



Random Policy



Optimal Policy

- The arrows indicate equal probability of moving into each of the directions

# Solving for the Optimal Policy

# Bellman Optimality Equation

- ▶ The **Bellman Optimality Equation** is named after Richard Ernest Bellman who introduced **dynamic programming** in 1953
- ▶ Almost any problem which can be solved using **optimal control theory** can be solved via the appropriate Bellman equation



Richard Ernest Bellman

# Bellman Optimality Equation

The **Bellman Optimality Equation (BOE)** decomposes  $Q^*$  as follows:

$$\begin{aligned} Q^*(s_t, a_t) &= \mathbb{E} [r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t, a_t] \\ &\stackrel{BOE}{=} \mathbb{E} \left[ r_t + \gamma \max_{a' \in \mathcal{A}} Q^*(s_{t+1}, a') \middle| s_t, a_t \right] \end{aligned}$$

This **recursive formulation** comprises two parts:

- ▶ Current reward:  $r_t$
- ▶ Discounted optimal action-value of successor:  $\gamma \max_{a' \in \mathcal{A}} Q^*(s_{t+1}, a')$

We want to **determine**  $Q^*(s_t, a_t)$ . How can we **solve** the BOE?

- ▶ The BOE is non-linear (because of max-operator)  $\Rightarrow$  no closed form solution
- ▶ Several iterative methods have been proposed, most popular: Q-Learning

# Proof of the Bellman Optimality Equation

**Proof** of the **Bellman Optimality Equation** for the **optimal action-value function**  $Q^*$ :

$$\begin{aligned}Q^*(s_t, a_t) &= \mathbb{E} [r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t, a_t] \\&= \mathbb{E} \left[ \sum_{k \geq 0} \gamma^k r_{t+k} | s_t, a_t \right] \\&= \mathbb{E} \left[ r_t + \gamma \sum_{k \geq 0} \gamma^k r_{t+k+1} | s_t, a_t \right] \\&= \mathbb{E} [r_t + \gamma V^*(s_{t+1}) | s_t, a_t] \\&= \mathbb{E} \left[ r_t + \gamma \max_{a'} Q^*(s_{t+1}, a') | s_t, a_t \right]\end{aligned}$$

# Bellman Optimality Equation

## Why is it useful to solve the BOE?

- ▶ A greedy policy which chooses the action that maximizes the optimal action-value function  $Q^*$  or the optimal state-value function  $V^*$  takes into account the reward consequences of all possible future behavior
- ▶ Via  $Q^*$  and  $V^*$  the optimal expected long-term return is turned into a quantity that is locally and immediately available for each state / state-action pair
- ▶ For  $V^*$ , a one-step-ahead search yields the optimal actions
- ▶  $Q^*$  effectively caches the results of all one-step-ahead searches



# Q-Learning

**Q-Learning:** Iteratively solve for  $Q^*$

$$Q^*(s_t, a_t) = \mathbb{E} \left[ r_t + \gamma \max_{a' \in \mathcal{A}} Q^*(s_{t+1}, a') \middle| s_t, a_t \right]$$

by constructing an **update sequence**  $Q_1, Q_2, \dots$  using learning rate  $\alpha$ :

$$\begin{aligned} Q_{i+1}(s_t, a_t) &\leftarrow (1 - \alpha)Q_i(s_t, a_t) + \alpha(r_t + \gamma \max_{a' \in \mathcal{A}} Q_i(s_{t+1}, a')) \\ &= Q_i(s_t, a_t) + \underbrace{\alpha(r_t + \gamma \max_{a' \in \mathcal{A}} Q_i(s_{t+1}, a') - Q_i(s_t, a_t))}_{\text{temporal difference (TD) error}} \end{aligned}$$

targetprediction

►  $Q_i$  will converge to  $Q^*$  as  $i \rightarrow \infty$       Note: policy  $\pi$  learned implicitly via Q table!

# Q-Learning

## Implementation:

- ▶ Initialize Q table and initial state  $s_0$  randomly
- ▶ Repeat:
  - ▶ Observe state  $s_t$ , choose action  $a_t$  according to  $\epsilon$ -greedy strategy  
(Q-Learning is “off-policy” as the updated policy is different from the behavior policy)
  - ▶ Observe reward  $r_t$  and next state  $s_{t+1}$
  - ▶ Compute TD error:  $r_t + \gamma \max_{a' \in \mathcal{A}} Q_i(s_{t+1}, a') - Q_i(s_t, a_t)$
  - ▶ Update Q table

What's the problem with using Q tables?

- ▶ **Scalability:** Tables don't scale to high dimensional state/action spaces (e.g., GO)
- ▶ **Solution:** Use a function approximator (neural network) to represent  $Q(s, a)$


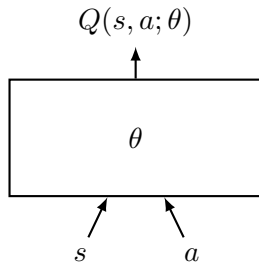
## 4.3

# Deep Q-Learning

# Deep Q-Learning

Use a **deep neural network** with weights  $\theta$  as function approximator to estimate  $Q$ :

$$Q(s, a; \theta) \approx Q^*(s, a)$$



# Training the Q Network

## Forward Pass:

Loss function is the mean-squared error in Q-values:

$$\mathcal{L}(\theta) = \mathbb{E} \left[ \left( \underbrace{r_t + \gamma \max_{a'} Q(s_{t+1}, a'; \theta)}_{\text{target}} - \underbrace{Q(s_t, a_t; \theta)}_{\text{prediction}} \right)^2 \right]$$

## Backward Pass:

Gradient update with respect to  $Q$ -function parameters  $\theta$ :

$$\nabla_{\theta} \mathcal{L}(\theta) = \nabla_{\theta} \mathbb{E} \left[ \left( r_t + \gamma \max_{a'} Q(s_{t+1}, a'; \theta) - Q(s_t, a_t; \theta) \right)^2 \right]$$

Optimize objective end-to-end with stochastic gradient descent (SGD) using  $\nabla_{\theta} \mathcal{L}(\theta)$ .

# Experience Replay

**To speed-up training** we like to train on **mini-batches**:

- ▶ Problem: Learning from consecutive samples is inefficient
- ▶ Reason: Strong correlations between consecutive samples

**Experience replay** stores agent's experiences at each time-step

- ▶ Continually update a **replay memory**  $D$  with new experiences  $e_t = (s_t, a_t, r_t, s_{t+1})$
- ▶ Train on samples  $(s_t, a_t, r_t, s_{t+1}) \sim U(D)$  drawn uniformly at random from  $D$
- ▶ Breaks correlations between samples
- ▶ Improves data efficiency as each sample can be used multiple times

In practice, a **circular replay memory** of finite memory size is used.

# Fixed Q Targets

## Problem: Non-stationary targets

- ▶ As the policy changes, so do our targets:  $r_t + \gamma \max_{a'} Q(s_{t+1}, a'; \theta)$
- ▶ This may lead to oscillation or divergence

## Solution: Use fixed Q targets to stabilize training

- ▶ A target network  $Q$  with weights  $\theta^-$  is used to generate the targets:

$$\mathcal{L}(\theta) = \mathbb{E}_{(s_t, a_t, r_t, s_{t+1}) \sim U(D)} \left[ \left( r_t + \gamma \max_{a'} Q(s_{t+1}, a'; \theta^-) - Q(s_t, a_t; \theta) \right)^2 \right]$$

- ▶ Target network  $Q$  is only updated every  $C$  steps by cloning the  $Q$ -network
- ▶ Effect: Reduces oscillation of the policy by adding a delay

# Putting it together

## Deep Q-Learning using experience replay and fixed Q targets:

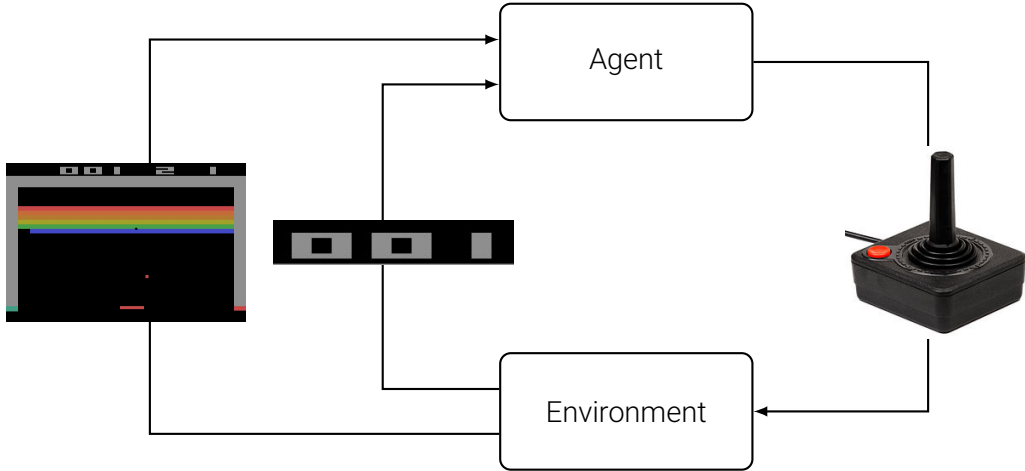
- ▶ Take action  $a_t$  according to  $\epsilon$ -greedy policy
- ▶ Store transition  $(s_t, a_t, r_t, s_{t+1})$  in replay memory  $D$
- ▶ Sample random mini-batch of transitions  $(s_t, a_t, r_t, s_{t+1})$  from  $D$
- ▶ Compute Q targets using old parameters  $\theta^-$
- ▶ Optimize MSE between Q targets and Q network predictions

$$\mathcal{L}(\theta) = \mathbb{E}_{s_t, a_t, r_t, s_{t+1} \sim D} \left[ \left( r_t + \gamma \max_{a'} Q(s_{t+1}, a'; \theta^-) - Q(s_t, a_t; \theta) \right)^2 \right]$$

using stochastic gradient descent.



# Case Study: Playing Atari Games



**Objective:** Complete the game with the highest score

# Case Study: Playing Atari Games

$Q(s, a; \theta)$ : Neural network with weights  $\theta$

FC-Out (Q values)

FC-256

32 4x4 conv, stride 2

16 8x8 conv, stride 2



**Output:** Q values for all (4 to 18) Atari actions  
(efficient: single forward pass computes Q for all actions)

**Input:**  $84 \times 84 \times 4$  stack of last 4 frames  
(after grayscale conversion, downsampling, cropping)

# Case Study: Playing Atari Games



# Deep Q-Learning Shortcomings

Deep Q-Learning suffers from several **shortcomings**:

- ▶ Long training times
- ▶ Uniform sampling from replay buffer  $\Rightarrow$  all transitions equally important
- ▶ Simplistic exploration strategy
- ▶ Action space is limited to a discrete set of actions  
(otherwise, expensive test-time optimization required)

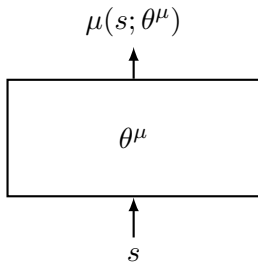
Various **improvements** over the original algorithm have been explored.

# Deep Deterministic Policy Gradients


**DDPG** addresses the problem of continuous action spaces.

**Problem:** Finding a continuous action requires optimization at every timestep.

**Solution:** Use two networks, an **actor** (deterministic policy) and a **critic**.



**Actor**



**Critic**

# Deep Deterministic Policy Gradients

- ▶ **Actor** network with weights  $\theta^\mu$  estimates agent's deterministic policy  $\mu(s; \theta^\mu)$ 
  - ▶ Update deterministic policy  $\mu(\cdot)$  in direction that most improves  $Q$
  - ▶ Apply chain rule to the **expected return** (this is the policy gradient):

$$\nabla_{\theta^\mu} \mathbb{E}_{s_t, a_t, r_t, s_{t+1} \sim D} [Q(s_t, \mu(s_t; \theta^\mu); \theta^Q)] = \mathbb{E} [\nabla_{a_t} Q(s_t, a_t; \theta^Q) \nabla_{\theta^\mu} \mu(s_t; \theta^\mu)]$$

- ▶ **Critic** estimates value of current policy  $Q(s, a; \theta^Q)$ 
  - ▶ Learned using the **Bellman Optimality Equation** as in Q Learning:

$$\nabla_{\theta^Q} \mathbb{E}_{s_t, a_t, r_t, s_{t+1} \sim D} \left[ (r_t + \gamma Q(s_{t+1}, \mu(s_{t+1}; \theta^{\mu-}); \theta^{Q-}) - Q(s_t, a_t; \theta^Q))^2 \right]$$

- ▶ Remark: No maximization over actions required as this step is now learned via  $\mu(\cdot)$

# Deep Deterministic Policy Gradients

**Experience replay** and **target networks** are again used to stabilize training:

- ▶ Replay memory  $D$  stores transition tuples  $(s_t, a_t, r_t, s_{t+1})$
- ▶ Target networks are updated using “soft” target updates
  - ▶ Weights are not directly copied but slowly adapted:

$$\begin{aligned}\theta^{Q-} &\leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q-} \\ \theta^{\mu-} &\leftarrow \tau \theta^{\mu} + (1 - \tau) \theta^{\mu-}\end{aligned}$$

where  $0 < \tau \ll 1$  controls the tradeoff between speed and stability of learning

**Exploration** is performed by adding noise  $\nabla_{\theta^{\mu}}$  to the policy  $\mu(s)$ :

$$\mu(s; \theta^{\mu}) + \mathcal{N}$$

# Prioritized Experience Replay

**Prioritize experience** to replay important transitions more frequently

- ▶ Priority  $\delta$  is measured by magnitude of temporal difference (TD) error:

$$\delta = |r_t + \gamma \max_{a'} Q(s_{t+1}, a'; \theta^{Q-}) - Q(s_t, a_t; \theta^Q)|$$

- ▶ TD error measures how “surprising” or unexpected the transition is
- ▶ Stochastic prioritization avoids overfitting due to lack of diversity
- ▶ Enables learning speed-up by a factor of 2 on Atari benchmarks



# Learning to Drive in a Day

## Real-world RL demo by Wayve:

- ▶ Deep Deterministic Policy Gradients with Prioritized Experience Replay
- ▶ **Input:** Single monocular image
- ▶ **Action:** Steering and speed
- ▶ **Reward:** Distance traveled without the safety driver taking control (requires no maps / localization)
- ▶ 4 Conv layers, 2 FC layers
- ▶ Only 35 training episodes

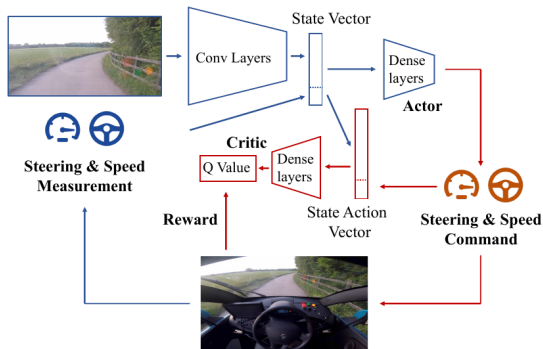


Fig. 1: We design a deep reinforcement learning algorithm for autonomous driving. This figure illustrates the actor-critic algorithm which we use to learn a policy and value function for driving. Our agent maximises the reward of distance travelled before intervention by a safety driver.

# Learning to Drive in a Day



## Other flavors of Deep RL

# Asynchronous Deep Reinforcement Learning


**Execute multiple agents** in separate environment instances:

- ▶ Each agent interacts with its own environment copy and collects experience
- ▶ Agents may use different exploration policies to maximize experience diversity
- ▶ Experience is not stored but directly used to update a shared global model
- ▶ Stabilizes training in similar way to experience replay by decorrelating samples
- ▶ Leads to reduction in training time roughly linear in the number of parallel agents

# Bootstrapped DQN

**Bootstrapping** for efficient exploration:

- ▶ Approximate a distribution over Q values via  $K$  bootstrapped "heads"
- ▶ At the start of each epoch, a single head  $Q_k$  is selected uniformly at random
- ▶ After training, all heads can be combined into a single ensemble policy



# Double Q-Learning

## Double Q-Learning

- Decouple Q function for selection and evaluation of actions to avoid Q overestimation and stabilize training. Target:

$$DQN \quad : \quad r_t + \gamma \max_{a'} Q(s_{t+1}, a'; \theta^-)$$

$$DoubleDQN \quad : \quad r_t + \gamma Q(s_{t+1}, \operatorname{argmax}_{a'} Q(s_{t+1}, a'; \theta); \theta^-)$$

- Online network with weights  $\theta$  is used to determine greedy policy
- Target network with weights  $\theta^-$  is used to determine corresponding action value
- Improves performance on Atari benchmarks

# Deep Recurrent Q-Learning

**Add recurrency** to a deep Q-network to handle *partial observability* of states:

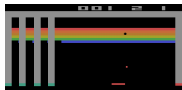
FC-Out (Q-values)

**LSTM**

Replace fully-connected layer with **recurrent LSTM layer**

32 4x4 conv, stride 2

16 8x8 conv, stride 2



# Faulty Reward Functions





# Summary

- ▶ Reinforcement learning learns through **interaction** with the environment
- ▶ The environment is typically modeled as a **Markov Decision Process**
- ▶ The goal of RL is to **maximize the expected future reward**
- ▶ Reinforcement learning requires trading off **exploration** and **exploitation**
- ▶ **Q-Learning** iteratively solves for the optimal action-value function
- ▶ The policy is learned implicitly via the **Q table**
- ▶ **Deep Q-Learning** scales to continuous/high-dimensional state spaces
- ▶ **Deep Deterministic Policy Gradients** scales to continuous action spaces
- ▶ Experience replay and target networks are necessary to stabilize training